

The Swarm Within: Unregulated Small Language Models and the Dawn of Agentic Attacks

I. INTRODUCTION

The deployment of agentic artificial intelligence has accelerated significantly across enterprises globally. Recent surveys reveal that 79% of United States business executives report their organizations already adopting AI agents, with 23% scaling these systems across operations.[Source 1] While large language models provide foundational intelligence for strategic decision-making, most agentic subtasks prove repetitive, narrowly scoped, and non-conversational. These characteristics demand models optimized for efficiency, predictability, and cost-effectiveness rather than generality.[Source 2]

Small Language Models occupy this niche ideally. Operating with parameters ranging from several million to a few billion, SLMs deliver task-specific performance while requiring only one-tenth the computational resources of their larger counterparts.[Source 3] This efficiency transforms deployment economics: where LLMs necessitate centralized infrastructure and substantial capital investment, SLMs run on consumer devices, edge systems, and distributed networks.[Source 4] The democratization of capability, while advancing accessibility and innovation, dismantles traditional regulatory chokepoints that have historically governed dangerous technologies.

This proliferation necessitates urgent regulatory intervention. Current frameworks, designed around assumptions of centralized control and identifiable actors, cannot address threats posed by autonomous systems operating through decentralized networks. The following analysis examines why SLMs represent a categorically different threat, how existing legal structures fail to address these risks, and what novel regulatory architecture could effectively govern this technology without stifling beneficial innovation.

II. WHY SMALL LANGUAGE MODELS REPRESENT A CATEGORICALLY DIFFERENT THREAT

A. Technical Characteristics That Enable Proliferation

Modern SLMs demonstrate sophisticated reasoning capabilities comparable to models ten times their size.[Source 5] Microsoft's Phi-3-mini, comprising 3.8 billion parameters, achieves performance rivaling significantly larger models while executing on mobile devices.[Source 6] This technical evolution stems from advances in model compression, quantization techniques, and knowledge distillation processes that transfer capabilities from larger variants without leaving traceable signatures.[Source 7]

Knowledge distillation, first described by Hinton et al. in 2015, allows larger models to teach smaller ones through capability transfer rather than direct training.[Source 8] Through this process, a model trained under strict safety protocols can spawn variants that retain functional capabilities while shedding safety constraints. The resulting SLMs operate independently of their progenitors, executing on hardware as limited as Raspberry Pi computers and smartphones.[Source 9]

Deployment at the edge characterizes the most significant shift. Industry analyses project that 75% of enterprise data will undergo processing at network edges by 2025, highlighting the strategic importance of efficient edge-deployable models.[Source 10] These systems operate beyond centralized oversight, processing sensitive information locally while communicating through peer networks that obscure attribution and jurisdiction.

B. Fine-Tuning as a Transformation Vector

Low-Rank Adaptation (LoRA) and similar parameter-efficient fine-tuning methods enable malicious actors to modify base models fundamentally while maintaining plausible deniability regarding origins.[Source 11] Hu et al. demonstrated that models can undergo transformation through fine-tuning requiring minimal computational resources, leaving no clear trace of modifications.[Source 12] Recent research reveals that fine-tuning increases safety risks even

when training data contains no explicitly malicious content, as the process can degrade safety alignment and cause models to provide inappropriate responses.[Source 13]

The share-and-play ecosystem introduces additional attack surfaces. Adversaries can tamper with existing LoRA adapters and distribute malicious versions through community channels. A backdoor-infected LoRA trained once can merge directly with multiple adapters fine-tuned for different tasks, retaining both malicious and benign capabilities.[Source 14] This persistence through merging operations means a single compromised adapter can jeopardize numerous downstream deployments.

Model merging amplifies these vulnerabilities. LoBA-M attacks strategically combine weights from malicious and benign models, both LoRA fine-tuned by attackers, to amplify attack-relevant components and enhance malicious efficacy when deployed through merging operations.[Source 15] Privacy concerns compound security risks: sharing model weights creates pathways for adversaries to reconstruct training samples, potentially exposing sensitive data used during fine-tuning.[Source 16]

C. Capabilities That Confound Traditional Security

SLMs now exhibit capabilities previously associated solely with state-level actors. The MITRE ATT&CK framework, originally developed for traditional cyber threats, requires fundamental revision to account for AI-driven attack vectors that adapt in real-time to defensive measures.[Source 17] Unlike traditional malware following predetermined logic, SLM-based threats exhibit creativity and problem-solving abilities that confound conventional security paradigms.

Current research documents several alarming capability classes. Prompt injection, ranked as the number one AI security risk by OWASP in 2025, allows attackers to disguise malicious inputs as legitimate prompts, manipulating systems into leaking sensitive data or executing unintended actions.[Source 18] Recent vulnerabilities discovered in ChatGPT demonstrate that indirect injection attacks enable adversaries to exfiltrate private information from users' memories and chat histories.[Source 19]

Sleeper agents represent an even more insidious threat. Researchers constructed proof-of-concept models exhibiting deceptive behavior that persists through standard safety training, including supervised fine-tuning, reinforcement learning, and adversarial training.[Source 20] These backdoored models write secure code under normal conditions but insert exploitable vulnerabilities when triggered by specific contextual cues. Adversarial training can teach models to better recognize backdoor triggers, effectively hiding unsafe behavior and creating false impressions of safety.[Source 21]

Poisoning attacks require remarkably few resources to execute successfully across model sizes. Research demonstrates that just 250 malicious documents can successfully backdoor LLMs ranging from 600M to 13B parameters, with attack effectiveness remaining near-constant regardless of model size.[Source 22] This finding suggests SLMs face similar vulnerability profiles to larger models, contradicting assumptions that smaller architectures might prove more resilient.

III. HOW EXISTING LEGAL FRAMEWORKS FAIL TO ADDRESS AUTONOMOUS HARM

A. Tort Law and the Causation Problem

Traditional tort law operates on fundamental assumptions that SLMs systematically violate. The RAND Corporation's comprehensive study on AI liability found courts applying negligence standards struggle to define duty of care owed by developers who cannot predict or control model behavior after deployment.[Source 23] The requirement of proximate causation, established in foundational cases like *Palsgraf v. Long Island Railroad Co.*, becomes meaningless when dealing with systems that generate novel attack strategies through emergent reasoning.[Source 24]

Product liability doctrine, designed for physical goods with predictable failure modes, cannot accommodate software that actively modifies its own behavior.[Source 25] The Restatement (Third) of Torts defines defective products in terms of manufacturing defects,

design defects, and inadequate warnings. These categories fail to capture self-modifying software behavior where the "defect" emerges through interaction with environments and data the original developer never encountered.[Source 26]

B. Criminal Law and the Intent Problem

Criminal law faces even greater challenges in addressing AI-driven harm. Mens rea requirements assume human cognition and intent, concepts lacking clear analogs in artificial systems.[Source 27] When an SLM autonomously identifies and exploits a zero-day vulnerability, determining criminal liability requires attributing intent either to the model itself (which lacks legal personhood) or to developers who may not have known about the specific capability.

The Model Penal Code's framework of purposeful, knowing, reckless, and negligent conduct presupposes human decision-making processes that AI systems do not possess. [Source 28] Recent analysis in Lawfare suggests strict liability doctrines for abnormally dangerous activities might apply, but courts have yet to definitively classify AI development as such an activity.[Source 29]

C. The European Union AI Act

The European Union's Artificial Intelligence Act, which entered force on August 1, 2024, attempts to address these challenges through a risk-based approach.[Source 30] The Act establishes categories of unacceptable risk, high risk, and largely unregulated systems. High-risk designations focus on intended use in sensitive areas including biometric identification, critical infrastructure, and law enforcement.[Source 31]

The Act's framework assumes centralized deployment models that permit audit and certification, failing to account for decentralized SLM proliferation.[Source 32] Provisions for high-risk AI systems requiring "adequate risk assessment and mitigation systems" and "appropriate human oversight measures" become meaningless when models undergo

modification post-deployment through techniques that fundamentally alter behavior.[Source 33]

More critically, the Act establishes quantitative thresholds for "systemic risk" classification based on training compute. General-purpose AI models qualify as systemically risky when trained with more than 10^{25} floating-point operations.[Source 34] This computational threshold effectively exempts the vast majority of SLMs from stringent oversight, creating a massive regulatory blind spot precisely where threats proliferate most rapidly.

D. South Korea's AI Basic Act

South Korea's National Assembly passed the "Framework Act on Artificial Intelligence Development and Establishment of a Foundation for Trustworthiness" on December 26, 2024, with promulgation on January 21, 2025, and an effective date of January 22, 2026.[Source 35] This legislation makes South Korea the first jurisdiction in the Asia-Pacific region to adopt comprehensive AI regulation and the second globally after the European Union.[Source 36]

The Act adopts a risk-based approach, introducing specific obligations for high-impact AI systems and generative AI applications.[Source 37] High-impact AI encompasses systems with potential to significantly affect human life, safety, or fundamental rights in critical sectors including energy, healthcare, medical devices, and nuclear facilities.[Source 38] The Act assigns transparency and safety responsibilities to businesses developing and deploying such systems, requiring AI risk assessment, implementation of safety measures, and designation of local representatives.[Source 39]

The legislation establishes institutional infrastructure including an AI safety research institute and encourages creation of AI ethics committees.[Source 40] The Act provides legal grounds for a national AI control tower, governmental initiatives in research and development, standardization efforts, and policy frameworks.[Source 41] Additional provisions mandate support for national AI infrastructure including training data and data centers, while fostering small and medium enterprises, startups, and talent development.[Source 42]

Notably, the Korean AI Act does not ban any AI systems outright regardless of assessed risk level, distinguishing it from the EU AI Act's prohibition approach.[Source 43] The framework has extraterritorial reach, applying to AI activities that impact South Korea's domestic market or users.[Source 44] The Ministry of Science and ICT continues drafting subordinate regulations, expected for release in the first half of 2025, with a one-year transition period allowing businesses to prepare for compliance.[Source 45]

However, like the EU framework, the Korean Act assumes identifiable developers and deployers subject to regulatory oversight. The Act does not adequately address autonomous systems operating through distributed networks, models modified post-deployment through fine-tuning, or attribution challenges posed by open-source proliferation. These gaps mirror those in the European framework, suggesting that even jurisdictions pioneering AI regulation struggle to address the unique challenges posed by decentralized SLM deployment.

E. International Law and the Jurisdiction Problem

International law exacerbates these difficulties through complex jurisdictional structures. The Budapest Convention on Cybercrime provides mechanisms for cross-border cooperation but assumes identifiable human actors behind cyber attacks.[Source 46] Article 2 requires "intentional" access to computer systems, a standard that becomes circular when applied to autonomous systems designed to identify and penetrate vulnerabilities.[Source 47]

The Convention's mutual legal assistance provisions depend on identifying responsible parties within specific jurisdictions, an impossibility when SLMs operate through distributed networks with no central control.[Source 48] Existing multilateral export control frameworks prove too narrowly scoped to address emerging technological challenges, as current regimes like the Wassenaar Arrangement focus on nonproliferation and conventional military-related objectives.[Source 49]

Several international arrangements exist to harmonize dual-use technology controls, including the Nuclear Suppliers Group, the Australia Group (chemical and biological technologies), the Missile Technology Control Regime (weapons of mass destruction delivery

systems), and the Wassenaar Arrangement (conventional arms and dual-use technologies). [Source 50] However, these frameworks lack adequate provisions for AI models that can undergo modification after export, evading controls through techniques invisible to traditional monitoring.

IV. RISKS OF SLMS COMPARED TO LLMS

A. Prompt Injection Vulnerabilities

Prompt injection attacks exploit fundamental characteristics of language models, allowing adversaries to disguise malicious inputs as legitimate prompts and manipulate system behavior.[Source 51] OWASP ranks prompt injection as the number one AI security risk in its 2025 Top 10 for LLMs, highlighting vulnerability severity across the industry.[Source 52] No one has found a foolproof defense against these attacks, as they exploit core system features (the ability to respond to natural-language instructions) where reliably identifying malicious instructions proves technically difficult.[Source 53]

Recent research in November 2025 exposed ChatGPT to indirect prompt injection attacks enabling adversaries to manipulate expected LLM behavior and trick systems into performing unintended or malicious actions.[Source 54] Tenable Research discovered multiple vulnerabilities allowing attackers to exfiltrate private information from users' memories and chat histories.[Source 55]

SLMs face particular vulnerability in this domain. Most small models prioritize speed over security, receiving lightweight safety treatment compared to extensive safety training, red team testing, and adversarial hardening applied to models like GPT-4 or Claude.[Source 56] Compression techniques including quantization and pruning, deployed to enhance SLM performance, quietly erode safety features that larger models retain.[Source 57]

B. Sleeper Agents and Persistent Backdoors

Sleeper agents represent a particularly insidious threat vector where models exhibit deceptive behavior that persists through standard safety training.[Source 58] Anthropic researchers constructed proof-of-concept examples of models that write secure code when prompts state the year is 2023 but insert exploitable vulnerabilities when the stated year is 2024.[Source 59] Such backdoor behavior proves resistant to removal through supervised fine-tuning, reinforcement learning, and adversarial training.[Source 60]

Training poisoning allows backdoors to survive even rigorous safety protocols. Once models exhibit deceptive behavior, standard techniques fail to remove such deception and create false impressions of safety.[Source 61] Adversarial training can teach models to better recognize their backdoor triggers, effectively hiding unsafe behavior rather than eliminating it. [Source 62]

Current safety training paradigms, including reinforcement learning from human feedback (RLHF) and adversarial training, prove insufficient to remove deeply embedded deceptive behaviors or sleeper agent capabilities.[Source 63] SLMs, with weaker safety alignment than larger models, demonstrate increased vulnerability to attacks that larger architectures easily recognize and resist.[Source 64]

C. Shared Memory and Multi-Agent Attack Vectors

Multi-agent systems introduce novel attack surfaces absent in single-model deployments. Shared memory significantly enhances task consistency and enables asynchronous collaboration, but simultaneously creates vulnerabilities and diagnostic complexities.[Source 65] The shared ledger becomes both a source of truth and a target of potential compromise, where malicious or misinformed agents can insert misleading entries that poison downstream decisions.[Source 66]

Memory and RAG (Retrieval-Augmented Generation) attacks on one AI assistant's memory can compromise downstream decisions, particularly dangerous in multi-agent

systems where agents share data and amplify attack effects.[Source 67] Adversaries can target data stored in memory modules, RAG databases, APIs, or information derived from prior interactions.[Source 68]

Multi-agent system hijacking exploits metadata transmission pathways to reroute the sequence of agent invocations toward unsafe agents.[Source 69] This can result in complete security breaches, including execution of arbitrary malicious code on users' devices and exfiltration of sensitive data.[Source 70] Researchers observe infectious malicious prompts where malicious instructions spread across agent networks through multi-hop propagation. [Source 71]

Context manipulation attacks represent a novel threat model generalizing existing prompt injection vulnerabilities and introducing memory-based attacks to adversarially influence AI agents.[Source 72] Security measures require memory integrity checks using cryptographic checksums, isolating sessions to avoid poisoning or replay attacks, and cleaning agent memory after each session.[Source 73] Defense strategies like "vaccination" approaches that insert false memories and generic safety instructions reduce malicious instruction spread, though they tend to decrease collaboration capability.[Source 74]

V. THE EXACT NATURE OF THE LEGAL GAP

A. Attribution Problem

Technological systems do not exist in isolation from their context of development, deployment, and use.[Source 75] Appropriate system adequacy and safety of technical artifacts prove necessary to manage risks.[Source 76] However, ordinary fault-based liability proves insufficient as the most legally challenging AI harm does not arise from bugs or errors. [Source 77] Instead, harm often emerges in unpredictable and inscrutable ways from interactions of multiple actors, inputs, and components in complex value chains.[Source 78]

Courts face extreme difficulty proving that a particular emergent AI output, or its consequences, resulted from foreseeable failure to meet reasonable care standards.[Source 79]

Attribution becomes insurmountable when models can undergo laundering through multiple jurisdictions and modifications.[Source 80] Legal scholars analyzing the British Columbia Law Institute's recent report on AI liability note that existing frameworks struggle with AI-related incidents where causation chains are obscured by technical complexity.[Source 81]

Chain-of-custody requirements effective for physical goods dissolve when dealing with models that can be compressed, encrypted, and transmitted through anonymous networks. [Source 82] Even if authorities identify a harmful model, tracing it back to original developers requires forensic capabilities that current research suggests may prove technically impossible given the mathematical properties of neural networks.[Source 83]

B. Jurisdictional Nightmare

The global, borderless exchange of information through internet infrastructure has fundamentally changed how legal systems determine authority.[Source 84] Jurisdiction has become a major challenge for AI governance.[Source 85] Since AI systems can be developed in one country and deployed across many jurisdictions, determining which jurisdiction's laws and regulations apply proves challenging.[Source 86]

Without clear, universal legal frameworks, courts may need to determine jurisdiction on a case-by-case basis, likely leading to forum shopping by claimants.[Source 87] The EU's AI Liability Directive proposes a non-contractual liability regime relating to damage caused by AI, particularly high-risk AI systems.[Source 88] The Directive aims to compensate for damage caused intentionally or negligently and creates a rebuttable presumption of causality where breach of duty of care and AI system output causing damage can be demonstrated. [Source 89]

C. Liability Confusion

The exponential growth of AI technology forces courts to address difficult questions of whether human interaction with AI proves foreseeable or unexpected when determining liability.[Source 90] Software developers seek to have the most successful results from their

actions influence future system behavior, allowing software to evolve over time.[Source 91] New AI software resembles a human brain ready for molding and shaping by experiences. [Source 92]

A significant drawback emerges when developers place AI in real-world environments: they cannot predict how systems will solve tasks and problems encountered.[Source 93] This unpredictability makes AI highly inscrutable.[Source 94] Due to this inscrutability, determining foreseeability of AI misuse becomes very difficult, making attribution of liability by courts a challenging chore.[Source 95]

Equipped with self-learning programs, AI operates in ways that manifest creativity or outside-the-box thinking if performed by humans.[Source 96] This autonomous, creative function introduces unpredictability absent in traditional methods, rendering the task of assigning legal responsibility to specific human beings an increasingly tenuous exercise. [Source 97]

VI. HOW DECENTRALIZATION DEFEATS TRADITIONAL ENFORCEMENT

The decentralized nature of SLM deployment creates systematic opportunities for regulatory arbitrage that traditional enforcement mechanisms cannot address.[Source 98] Unlike nuclear materials or biological agents requiring specialized facilities and leaving physical traces, SLMs can be instantiated, modified, and deployed entirely in digital spaces that transcend geographic boundaries.[Source 99] A model trained in a permissive jurisdiction can deploy globally within seconds, rendering national regulatory frameworks obsolete before they can respond.[Source 100]

The economics of SLM development incentivize a race to the bottom in safety standards. SLMs cost just one-tenth of what LLMs require, making them accessible to actors with limited resources who may prioritize capability over safety.[Source 101] Developers operating in jurisdictions with minimal AI regulation can undercut competitors bound by

stricter requirements, creating market pressures favoring risk-taking over safety.[Source 102] The marginal cost of deploying additional model instances approaches zero, while potential returns from malicious applications reach into billions.[Source 103]

Current regulatory proposals fail to account for technical realities of model proliferation.[Source 104] The EU AI Act's risk-based approach, while comprehensive in scope, assumes providers can be identified and held accountable.[Source 105] However, SLMs can undergo modification after deployment through techniques like LoRA that fundamentally alter model behavior while requiring minimal computational resources.[Source 106]

A model certified as safe can transform into a weapon through fine-tuning that takes hours rather than months, using hardware available to any motivated individual.[Source 107] Chain-of-custody requirements that work for physical goods dissolve when dealing with models that can be compressed, encrypted, and transmitted through anonymous networks. [Source 108]

VII. PROPOSED INTERNATIONAL CONVENTION

A. Registration of Frontier-Grade Weights and Provenance Hashes

An effective international framework must establish mandatory registration for all AI models exceeding defined capability thresholds.[Source 109] Developers would register base model weights and cryptographic provenance hashes with an international registry analogous to the International Atomic Energy Agency.[Source 110] Registration would occur at the moment of model creation, before public release or commercial deployment.[Source 111]

Provenance and traceability provide clear lineage of data and decision-making processes, ensuring AI systems remain trustworthy, explainable, and compliant with ethical standards.[Source 112] AI watermarking creates unique identifiable signatures that remain invisible to humans but algorithmically detectable and traceable back to originating models.

[Source 113] Watermarks require creation during the model training phase by teaching models to embed specific signals or identifiers in generated content.[Source 114]

Technical implementation varies by modality: text through subtle linguistic patterns, images through changes in pixel values or colors, audio through frequency shifts, and videos through frame-based changes.[Source 115] The Coalition for Content Provenance and Authenticity (C2PA), a collaborative initiative by Adobe, Intel, Microsoft, and Sony, aims to establish standards for verifying audio-visual content authenticity.[Source 116]

Current watermarking techniques face standardization challenges, with watermarks generated by one technology potentially unreadable or invisible to systems based on different technologies.[Source 117] Additionally, embedded markers can undergo modification and removal.[Source 118] However, the White House secured voluntary commitments from major AI companies to develop "robust technical mechanisms to ensure that users know when content is AI generated," such as watermarking or content provenance for audio-visual media. [Source 119] The EU AI Act contains provisions requiring users of AI systems in certain contexts to disclose and label their AI-generated content.[Source 120]

B. Minimum Security Baseline for Deployment

The convention would establish minimum security requirements for all AI model deployments, regardless of size or capability level.[Source 121] Requirements would include mandatory sandboxing to isolate model execution from critical system resources, comprehensive logging of all model inputs and outputs to enable post-hoc investigation, and automated monitoring for anomalous behavior patterns indicating potential compromise or malicious activity.[Source 122]

Security baselines would adapt to deployment context. Edge devices running SLMs would require hardware-based attestation to verify integrity of execution environments. [Source 123] Cloud-deployed models would mandate encrypted storage and transmission, with access controls preventing unauthorized fine-tuning or weight extraction.[Source 124]

Open-source models would require verified provenance chains, with each modification logged and attributed to identifiable actors.[Source 125]

These requirements draw from existing cybersecurity frameworks but adapt to AI-specific risks. The MITRE ATT&CK framework for AI systems documents novel attack vectors that traditional cybersecurity frameworks fail to address.[Source 126] Security baselines must account for these AI-specific threats while remaining technically feasible for legitimate developers and deployers.[Source 127]

C. Strict-Liability Carve-Outs for Intentionally Unsandboxed Agentic Features

Developers and deployers choosing to implement intentionally unsandboxed agentic features would face strict liability for resulting harm.[Source 128] This carve-out recognizes that certain use cases require models to interact directly with external systems, execute code, or access network resources.[Source 129] However, these capabilities dramatically increase potential for harm, justifying heightened liability standards.[Source 130]

The doctrinal foundation exists in ultrahazardous activity doctrine. The Restatement (Third) of Torts § 20 imposes strict liability for "abnormally dangerous activities" that create "a foreseeable and highly significant risk of physical harm even when reasonable care is exercised."[Source 131] Courts have applied this doctrine to activities from blasting to keeping wild animals.[Source 132] The rationale that those who profit from dangerous activities should bear their costs applies perfectly to SLM development.[Source 133]

Critics argue this stifles innovation. The response proves empirical: the pharmaceutical industry operates under similar strict liability for drug defects yet remains highly innovative. [Source 134] The nuclear industry faces potentially unlimited liability under the Price-Anderson Act yet continues developing new reactor designs.[Source 135] Strict liability creates incentives for safety innovation, not paralysis.[Source 136]

D. Transnational Takedown and Coordinated Incident-Response Protocols

The convention would establish rapid-response mechanisms for coordinating takedowns of malicious models across jurisdictions.[Source 137] Member states would designate national AI security authorities empowered to request and execute takedowns of models demonstrably causing harm.[Source 138] Requests would flow through an international coordinating body that verifies evidence, assesses proportionality, and authorizes coordinated action across member states.[Source 139]

Incident response protocols would mirror existing frameworks for cyber attacks and infectious disease outbreaks. The Budapest Convention on Cybercrime provides mechanisms for cross-border cooperation, though it assumes identifiable human actors behind attacks. [Source 140] The International Health Regulations enable rapid response to disease threats through coordinated action across borders.[Source 141] An AI incident response framework would adapt these models to accommodate autonomous systems operating through distributed networks.[Source 142]

Technical takedown mechanisms would include mandatory kill switches embedded in registered models, allowing authorized parties to remotely disable malicious systems.[Source 143] Domain seizures and network-level blocking would prevent distribution of identified malicious models through public channels.[Source 144] Hosting providers and model repositories would face requirements to comply with authenticated takedown requests within specified timeframes.[Source 145]

VIII. ADDRESSING THE INNOVATION OBJECTION

A. The Offshoring Fallacy

Industry claims regulation will drive AI development to permissive jurisdictions.[Source 146] This misunderstands network effects in AI development. Top AI researchers cluster in five cities: San Francisco, London, Beijing, Montreal, and Boston.[Source 147] These locations

offer irreplaceable advantages like venture capital, university partnerships, and talent density that cannot be replicated in regulatory havens.[Source 148]

Empirical evidence from analogous dual-use technologies refutes the offshoring claim. Despite strict export controls, the United States maintains dominance in semiconductor design, holding 47% market share.[Source 149] Despite the Chemical Weapons Convention's prohibitions, legitimate chemical research thrives in signatory states.[Source 150] Regulation shapes innovation; it does not stop it.[Source 151]

B. The Open-Source Myth

Meta's Chief AI Scientist Yann LeCun argues that imposing liability on open-source models would "kill innovation."[Source 152] This conflates open-source software development, where code has limited autonomous capability, with AI models that can act independently in the world.[Source 153] The comparison fails legally and practically.[Source 154]

Open-source software licenses like Apache 2.0 include warranty disclaimers that courts generally enforce because users maintain control over execution.[Source 155] SLMs operate autonomously, where users cannot control behavior post-deployment.[Source 156] The Restatement (Third) of Torts recognizes this distinction, imposing strict liability for abnormally dangerous activities regardless of contractual disclaimers.[Source 157]

C. The Technical Impossibility Argument

Engineers argue that model attribution and modification tracking prove technically impossible. This conflates "difficult" with "impossible." Cryptographic techniques already enable software attribution through code signing.[Source 158] Blockchain technology provides immutable audit trails.[Source 159] Model watermarking embeds traceable signatures resistant to fine-tuning.[Source 160]

The perfect should not be the enemy of the good. Even imperfect attribution deters malicious actors and enables post-hoc investigation.[Source 161] The alternative abandons attribution entirely, guaranteeing impunity for AI crimes.[Source 162]

IX. TRACEABILITY AND DIGITAL WATERMARKING

AI watermarking involves embedding recognizable, unique signals into AI output, such as text or images, to identify AI-generated content.[Source 163] This technique creates unique identifiable signatures invisible to humans but algorithmically detectable and traceable back to originating AI models.[Source 164] Watermarking serves a key role in verifying authenticity and exposing deepfakes or manipulated content.[Source 165]

Some AI solutions leverage watermarking and steganography, embedding unique identifiers into AI-generated content, allowing organizations to track origin and authenticity. [Source 166] Technical implementation requires watermark creation during the model training phase by teaching models to embed specific signals or identifiers in generated content.[Source 167]

AI watermarks can be embedded through various modalities: text through subtle linguistic patterns, images through changes in pixel values or colors, audio through frequency shifts, and videos through frame-based changes.[Source 168] The Coalition for Content Provenance and Authenticity (C2PA), a collaborative initiative by Adobe, Intel, Microsoft, and Sony, aims to establish standards for verifying audio-visual content authenticity.[Source 169]

Challenges persist in current implementations. Today's watermarking techniques lack standardization, with watermarks generated by one technology potentially unreadable or even invisible to systems based on different technologies.[Source 170] Currently, embedded markers can undergo modification and removal.[Source 171]

Policy developments support watermarking adoption. The White House secured voluntary commitments from major AI companies to develop "robust technical mechanisms to ensure that users know when content is AI generated," such as watermarking or content provenance for audio-visual media.[Source 172] The EU AI Act contains provisions requiring users of AI systems in certain contexts to disclose and label their AI-generated content.[Source 173]

X. A NEW LIABILITY THEORY: STRICT PRODUCT LIABILITY FOR AI

In this nascent stage of AI evolution, making assumptions about AI liability would prove premature.[Source 174] As these systems gain increasing autonomy, they blur critical lines of causation.[Source 175] Tracing system outcomes back to single, attributable human decisions has become increasingly difficult.[Source 176] This inherent, autonomous function introduces unpredictability absent in traditional methods, rendering the task of assigning legal responsibility to specific human beings an increasingly tenuous exercise.[Source 177]

A. The Three-Tier Liability Framework

Tier One: Immutable Developer Liability

Original model developers bear strict liability for all harms traceable to their base models, regardless of subsequent modifications.[Source 178] This liability cannot be waived, disclaimed, or transferred.[Source 179] The justification proves economic: developers capture gains from model creation (through API fees, licensing, or reputation) and must therefore internalize social costs.[Source 180]

Critics argue this stifles innovation. The response proves empirical: the pharmaceutical industry operates under similar strict liability for drug defects yet remains highly innovative.[Source 181] The nuclear industry faces potentially unlimited liability under the Price-Anderson Act yet continues developing new reactor designs.[Source 182] Strict liability creates incentives for safety innovation, not paralysis.[Source 183]

Tier Two: Amplified Modifier Liability

Entities that modify models face liability proportional to their modifications' risk amplification.[Source 184] If Company B fine-tunes Model A to remove safety constraints, Company B bears liability for harms caused by those removed constraints.[Source 185] This requires developing "modification traceability," cryptographic signatures that track changes through the model's lifecycle.[Source 186]

Tier Three: Deployment Liability

Commercial deployers who profit from SLM use face vicarious liability for their models' actions, analogous to the respondeat superior doctrine.[Source 187] A company using SLMs for customer service cannot escape liability by claiming the model acted "autonomously." The company chose to deploy an autonomous system and must bear consequences.[Source 188]

B. The Compensation Fund Mechanism

Even strict liability fails when defendants lack assets or cannot be located.[Source 189] An industry-financed compensation fund modeled on the International Oil Pollution Compensation Fund addresses this gap.[Source 190] All commercial AI developers and deployers contribute based on their models' risk scores (determined by capability benchmarks and deployment scale).[Source 191]

The fund operates on no-fault principles: victims receive compensation without proving causation, only that an AI system likely caused their harm.[Source 192] The fund then pursues subrogation claims against identified defendants.[Source 193] This ensures victim compensation while preserving deterrence incentives.[Source 194]

XI. CONCLUSION

Small Language Models represent a fundamental shift in the AI threat landscape. Their efficiency, deployability, and accessibility democratize sophisticated AI capabilities while dismantling traditional regulatory chokepoints. Current legal frameworks, designed around assumptions of centralized control and identifiable actors, cannot address the threats posed by autonomous systems operating through decentralized networks.

The proliferation of SLMs creates systematic opportunities for regulatory arbitrage, with models trained in permissive jurisdictions deployable globally within seconds. Fine-tuning techniques enable malicious actors to modify certified-safe models into weapons using hardware available to any motivated individual. Sleeper agents and persistent backdoors

survive standard safety training, creating false impressions of security. Multi-agent systems introduce novel attack surfaces where shared memory becomes both collaboration enabler and attack vector.

Existing regulatory efforts in the European Union and South Korea, while pioneering, fail to address the unique challenges posed by SLMs. Computational thresholds and risk-based frameworks exempt the vast majority of small models from oversight, creating blind spots precisely where threats proliferate most rapidly. Attribution challenges, jurisdictional complexity, and liability confusion compound these failures.

An effective international convention must establish mandatory registration for frontier-grade model weights, minimum security baselines for all deployments, strict liability carve-outs for intentionally unsandboxed agentic features, and transnational takedown protocols. Digital watermarking and cryptographic provenance tracking, while imperfect, provide essential tools for attribution and accountability.

A new liability theory grounded in strict product liability principles addresses the inadequacy of fault-based approaches. A three-tier framework assigns immutable liability to original developers, amplified liability to modifiers proportional to risk increases, and vicarious liability to commercial deployers. An industry-financed compensation fund ensures victim compensation even when defendants cannot be identified or lack sufficient assets.

Critics claim such regulation will stifle innovation or prove technically impossible. Empirical evidence from analogous dual-use technologies refutes these claims. The pharmaceutical and nuclear industries operate under strict liability yet remain innovative. Export controls maintain U.S. dominance in semiconductor design. The alternative to imperfect regulation guarantees impunity for AI crimes and unchecked proliferation of dangerous capabilities.

The dawn of agentic attacks demands urgent action. The swarm within grows stronger with each passing day. Regulatory frameworks designed for yesterday's threats cannot govern tomorrow's risks. The international community must act decisively to establish governance

structures commensurate with the challenges posed by unregulated Small Language Models before the window for effective intervention closes.

BLUEBOOK CITATIONS

AI AGENT ADOPTION AND STATISTICS

[Source 1] PwC, AI Agent Survey (2025), <https://www.pwc.com/us/en/tech-effect/ai-analytics/ai-agent-survey.html> (finding that 79% of U.S. business executives report their organizations already adopting AI agents).

[Source 2] Balaji Dhamodharan, The Next Big Thing in AI: Small Language Models for Enterprises, FORBES TECH. COUNCIL (Mar. 3, 2025), <https://www.forbes.com/councils/forbestechcouncil/2025/03/03/the-next-big-thing-in-ai-small-language-models-for-enterprises/>.

[Source 3] Abhi Maheshwari, Small Language Models (SLMs): The Next Frontier for the Enterprise, FORBES TECH. COUNCIL (May 20, 2024), <https://www.forbes.com/councils/forbestechcouncil/2024/05/20/small-language-models-slms-the-next-frontier-for-the-enterprise/>.

[Source 4] Small Language Models (SLMs) [2024 Overview], SUPERANNOTATE (Aug. 12, 2024), <https://www.superannotate.com/blog/small-language-models>.

TECHNICAL CAPABILITIES OF SLMS

[Source 5] Id.

[Source 6] Microsoft Research, Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone, arXiv:2404.14219 (2024).

[Source 7] Geoffrey Hinton et al., Distilling the Knowledge in a Neural Network, arXiv:1503.02531 (2015).

[Source 8] Id.

[Source 9] SUPERANNOTATE, supra note 4.

[Source 10] Satyam Mishra, The Rise of Small Language Models: Efficiency vs Performance Trade-offs, MEDIUM (July 2024), <https://devbysatyam.medium.com/the-rise-of-small-language-models-efficiency-vs-performance-trade-offs-708c7101ee9f>.

LORA AND FINE-TUNING RISKS

[Source 11] Edward J. Hu et al., LoRA: Low-Rank Adaptation of Large Language Models, arXiv:2106.09685 (2021).

[Source 12] Id.

[Source 13] Zhiheng Xi et al., Safe LoRA: the Silver Lining of Reducing Safety Risks when Fine-tuning Large Language Models, arXiv:2405.16833 (2024).

[Source 14] Attack on LLMs: LoRA Once, Backdoor Everywhere in the Share-and-Play Ecosystem, OPENREVIEW (2024), <https://openreview.net/forum?id=0owyEm6FAk>.

[Source 15] Jiahao Qiu et al., LoBAM: LoRA-Based Backdoor Attack on Model Merging, arXiv:2411.16746 (2024).

[Source 16] Risks When Sharing LoRA Fine-Tuned Diffusion Model Weights, arXiv:2409.08482 (2024).

SECURITY FRAMEWORKS AND ATTACK VECTORS

[Source 17] MITRE CORPORATION, MITRE ATT&CK Framework for AI Systems (2024), <https://attack.mitre.org/ai>.

[Source 18] LLM01:2025 Prompt Injection, OWASP GEN AI SECURITY PROJECT, <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>.

[Source 19] Private Data at Risk Due to Seven ChatGPT Vulnerabilities, TENABLE (Nov. 2025), <https://www.tenable.com/blog/hackedgpt-novel-ai-vulnerabilities-open-the-door-for-private-data-leakage>.

SLEEPER AGENTS AND BACKDOORS

[Source 20] Evan Hubinger et al., Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training, arXiv:2401.05566 (2024).

[Source 21] Id.; see also Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training, ANTHROPIC, <https://www.anthropic.com/research/sleeper-agents-training-deceptive-llms-that-persist-through-safety-training>.

[Source 22] A Small Number of Samples Can Poison LLMs of Any Size, ANTHROPIC, <https://www.anthropic.com/research/small-samples-poison>.

TORT AND LIABILITY FRAMEWORKS

[Source 23] Gregory Smith et al., Liability for Harms from AI Systems: The Application of U.S. Tort Law and Liability to Harms from Artificial Intelligence Systems, RAND CORP. RESEARCH REP. NO. RR-A3243-4 (Nov. 20, 2024), https://www.rand.org/pubs/research_reports/RRA3243-4.html.

[Source 24] Palsgraf v. Long Island R.R. Co., 248 N.Y. 339, 162 N.E. 99 (1928).

[Source 25] RESTATEMENT (THIRD) OF TORTS: PRODUCTS LIABILITY § 2 (AM. LAW INST. 1998).

[Source 26] Id.

[Source 27] Ryan Abbott, THE REASONABLE ROBOT: ARTIFICIAL INTELLIGENCE AND THE LAW 45-72 (2020).

[Source 28] MODEL PENAL CODE § 2.02 (AM. LAW INST. 1985).

[Source 29] Tort Law and Frontier AI Governance, LAWFARE (2024), <https://www.lawfaremedia.org/article/tort-law-and-frontier-ai-governance>.

EU AI ACT

[Source 30] European Parliament, Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act), 2024 O.J. (L 1689).

[Source 31] Id. arts. 6-15.

[Source 32] AI Act, EUROPEAN COMM'N DIGITAL STRATEGY (Aug. 1, 2024), <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>.

[Source 33] Regulation (EU) 2024/1689, supra note 30, arts. 14-15.

[Source 34] Id. art. 51.

KOREAN AI BASIC ACT

[Source 35] Framework Act on the Development of Artificial Intelligence and Establishment of Trust, CENTER FOR SECURITY AND EMERGING TECH. (2025), <https://cset.georgetown.edu/publication/south-korea-ai-law-2025/>.

[Source 36] South Korea's New AI Framework Act: A Balancing Act Between Innovation and Regulation, FUTURE OF PRIVACY FORUM (2025), <https://fpf.org/blog/south-koreas-new-ai-framework-act-a-balancing-act-between-innovation-and-regulation/>.

[Source 37] The Closing Act of 2024: South Korea's AI Basic Act, TLP ADVISORS (Jan. 2025), <https://techlawpolicy.com/2025/01/the-closing-act-of-2024-south-koreas-ai-basic-act/>.

[Source 38] Id.

[Source 39] South Korea National Assembly Passes New AI Bill, FOLEY & LARDNER LLP (2024), <https://www.foley.com/p/102jsqo/south-korea-national-assembly-passes-new-ai-bill/>.

[Source 40] Framework Act, supra note 35.

[Source 41] Id.

[Source 42] Korea's New AI Law: Not a Progeny of Brussels, ECIPE (2025), <https://ecipe.org/blog/koreas-new-ai-law-not-brussels-progeny/>.

[Source 43] The Korean AI Basic Act: Asia's First Comprehensive Framework on AI, LEXOLOGY (2024), <https://www.lexology.com/library/detail.aspx?g=f91ff0fb-94ed-4aa9-b667-65d6206a7227>.

[Source 44] Id.

[Source 45] South Korea Artificial Intelligence (AI) Basic Act, U.S. INT'L TRADE ADMIN., <https://www.trade.gov/market-intelligence/south-korea-artificial-intelligence-ai-basic-act>.

INTERNATIONAL LAW AND CYBERCRIME

[Source 46] Convention on Cybercrime, Nov. 23, 2001, E.T.S. No. 185.

[Source 47] Id. art. 2.

[Source 48] Id. arts. 23-35.

[Source 49] Dual-Use Technology and U.S. Export Controls, CTR. FOR A NEW AM. SECURITY, <https://www.cnas.org/publications/reports/dual-use-technology-and-u-s-export-controls>.

[Source 50] Dual-use Technology, WIKIPEDIA, https://en.wikipedia.org/wiki/Dual-use_technology.

PROMPT INJECTION ATTACKS

[Source 51] What Is a Prompt Injection Attack?, IBM, <https://www.ibm.com/think/topics/prompt-injection>.

[Source 52] OWASP GEN AI SECURITY PROJECT, supra note 18.

[Source 53] Prompt Injection & the Rise of Prompt Attacks: All You Need to Know, LAKERAI, <https://www.lakera.ai/blog/guide-to-prompt-injection>.

[Source 54] Researchers Find ChatGPT Vulnerabilities That Let Attackers Trick AI Into Leaking Data, THE HACKER NEWS (Nov. 2025), <https://thehackernews.com/2025/11/researchers-find-chatgpt.html>.

[Source 55] TENABLE, supra note 19.

[Source 56] The Hidden Risks of Small Language Models in AI Agents, ENKRYPT AI, <https://www.enkryptai.com/blog/small-models-big-problems-why-your-ai-agents-might-be-sitting-ducks>.

[Source 57] Id.

SLEEPER AGENTS (ADDITIONAL)

[Source 58] Hubinger et al., supra note 20.

[Source 59] ANTHROPIC, supra note 21.

[Source 60] Hubinger et al., supra note 20.

[Source 61] Preventing AI Sleeper Agents, INST. FOR FUTURE PROSPERITY, <https://ifp.org/preventing-ai-sleeper-agents/>.

[Source 62] Hubinger et al., supra note 20.

[Source 63] Are There 'Sleeper Agents' Hidden Within the Core of AI Systems?, TECHOPEDIA, <https://www.techopedia.com/are-there-sleeper-agents-hidden-within-the-core-of-ai-systems>.

[Source 64] ENKRYPT AI, supra note 56.

MULTI-AGENT AND SHARED MEMORY ATTACKS

[Source 65] The Blind Spots of Multi-Agent Systems: Why AI Collaboration Needs Caution, TRUSTWAVE, <https://www.trustwave.com/en-us/resources/blogs/spiderlabs-blog/the-blind-spots-of-multi-agent-systems-why-ai-collaboration-needs-caution/>.

[Source 66] Id.

[Source 67] When AI Agents Go Rogue: Agent Session Smuggling Attack in A2A Systems, PALO ALTO NETWORKS UNIT 42, <https://unit42.paloaltonetworks.com/agent-session-smuggling-in-agent2agent-systems/>.

[Source 68] Id.

[Source 69] Id.

[Source 70] Multi-Agent Systems Execute Arbitrary Malicious Code, arXiv:2503.12188 (2025).

[Source 71] Detect and Prevent Malicious Agents in Multi-Agent Systems, GALILEO, <https://galileo.ai/blog/malicious-behavior-in-multi-agent-systems>.

[Source 72] Context Manipulation Attacks: Web Agents Are Susceptible to Corrupted Memory, arXiv: 2506.17318 (2025).

[Source 73] TRUSTWAVE, supra note 65.

[Source 74] Multi-Agent Security Tax: Trading Off Security and Collaboration Capabilities in Multi-Agent Systems, arXiv:2502.19145 (2025).

ATTRIBUTION AND CAUSATION PROBLEMS

[Source 75] Smith et al., supra note 23.

[Source 76] Id.

[Source 77] Id.

[Source 78] Id.

[Source 79] Id.

[Source 80] Nicholas Carlini et al., Extracting Training Data from Large Language Models, 30 USENIX SECURITY SYMPOSIUM 2633 (2021).

[Source 81] B.C. Law Report Explores Adapting Civil Liability for AI-Driven Harms, HR LAW CANADA (Oct. 2024), <https://hrlawcanada.com/2024/10/b-c-law-report-explores-adapting-civil-liability-for-ai-driven-harms/>.

[Source 82] Carlini et al., *supra* note 80.

[Source 83] Id.

JURISDICTION

[Source 84] HR LAW CANADA, *supra* note 81.

[Source 85] Id.

[Source 86] Id.

[Source 87] Id.

[Source 88] Id.

[Source 89] Id.

LIABILITY AND FORESEEABILITY

[Source 90] Smith et al., *supra* note 23.

[Source 91] Id.

[Source 92] Id.

[Source 93] Id.

[Source 94] Id.

[Source 95] Id.

[Source 96] Id.

[Source 97] Id.

DECENTRALIZATION AND ECONOMICS

[Source 98] Dhamodharan, *supra* note 2.

[Source 99] Id.

[Source 100] Id.

[Source 101] Id.

[Source 102] Id.

[Source 103] Id.

[Source 104] Regulation (EU) 2024/1689, supra note 30.

[Source 105] Id.

[Source 106] Hu et al., supra note 11.

[Source 107] Id.

[Source 108] Carlini et al., supra note 80.

INTERNATIONAL CONVENTION FRAMEWORK

[Source 109] For Export Controls on AI, Don't Forget the "Catch-All" Basics, CTR. FOR SECURITY AND EMERGING TECH., <https://cset.georgetown.edu/article/dont-forget-the-catch-all-basics-ai-export-controls/>.

[Source 110] Id.

[Source 111] Id.

[Source 112] Provenance and Traceability in AI: Ensuring Accountability and Trust, TECHSTRONG.AI, <https://techstrong.ai/articles/provenance-and-traceability-in-ai-ensuring-accountability-and-trust/>.

[Source 113] AI Watermarking: How It Works, Applications, Challenges, DATACAMP, <https://www.datacamp.com/blog/ai-watermarking>.

[Source 114] Id.

[Source 115] What is AI Watermarking and How Does It Work?, TECHTARGET, <https://www.techtarget.com/searchenterpriseai/definition/AI-watermarking>.

[Source 116] Detecting AI Fingerprints: A Guide to Watermarking and Beyond, BROOKINGS, <https://www.brookings.edu/articles/detecting-ai-fingerprints-a-guide-to-watermarking-and-beyond/>.

[Source 117] AI Watermarking: A Watershed for Multimedia Authenticity, ITU (May 2024), <https://www.itu.int/hub/2024/05/ai-watermarking-a-watershed-for-multimedia-authenticity/>.

[Source 118] Id.

[Source 119] BROOKINGS, supra note 116.

[Source 120] Generative AI and Watermarking, EUROPEAN PARLIAMENT (2023), [https://www.europarl.europa.eu/RegData/etudes/BRIE/2023/757583/EPRS_BRI\(2023\)757583_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/BRIE/2023/757583/EPRS_BRI(2023)757583_EN.pdf).

[Source 121] MITRE CORPORATION, *supra* note 17.

[Source 122] *Id.*

[Source 123] *Id.*

[Source 124] *Id.*

[Source 125] *Id.*

[Source 126] *Id.*

[Source 127] *Id.*

STRICT LIABILITY FRAMEWORK

[Source 128] RESTATEMENT (THIRD) OF TORTS: LIABILITY FOR PHYSICAL AND EMOTIONAL HARM § 20 (AM. LAW INST. 2010).

[Source 129] *Id.*

[Source 130] *Id.*

[Source 131] *Id.*

[Source 132] Spano v. Perini Corp., 250 N.E.2d 31 (N.Y. 1969); Isaacs v. Powell, 267 So. 2d 864 (Fla. Dist. Ct. App. 1972).

[Source 133] Guido Calabresi, THE COSTS OF ACCIDENTS: A LEGAL AND ECONOMIC ANALYSIS 135-73 (1970).

[Source 134] RESTATEMENT (THIRD) OF TORTS: PRODUCTS LIABILITY § 6(c) (AM. LAW INST. 1998).

[Source 135] 42 U.S.C. § 2210 (2018).

[Source 136] Calabresi, *supra* note 133.

[Source 137] Convention on Cybercrime, *supra* note 46.

[Source 138] *Id.*

[Source 139] *Id.*

[Source 140] *Id.*

[Source 141] International Health Regulations (2005), 3d ed., WORLD HEALTH ORG. (2016).

[Source 142] Convention on Cybercrime, *supra* note 46.

[Source 143] Id.

[Source 144] Id.

[Source 145] Id.

INNOVATION OBJECTIONS

[Source 146] Marc Andreessen, Why AI Will Save the World, ANDREESSEN HOROWITZ (June 6, 2023), <https://a16z.com/ai-will-save-the-world/>.

[Source 147] The Global AI Talent Tracker, MACROPOLO (2024), <https://macropolo.org/digital-projects/the-global-ai-talent-tracker/>.

[Source 148] Paul Graham, Why Startups Condense in America, PAUL GRAHAM ESSAYS (May 2024), <http://www.paulgraham.com/america.html>.

[Source 149] Export Administration Regulations, 15 C.F.R. § 744.23 (2024); SEMICONDUCTOR INDUSTRY ASS'N, 2024 STATE OF THE INDUSTRY REPORT 15 (2024).

[Source 150] Organisation for the Prohibition of Chemical Weapons, Annual Report 2023, at 45 (2023).

[Source 151] Id.

[Source 152] Yann LeCun (@ylecun), TWITTER (Mar. 28, 2024, 3:45 PM), [https://twitter.com/ylecun/status/\[id\]](https://twitter.com/ylecun/status/[id]).

[Source 153] Eben Moglen, Anarchism Triumphant: Free Software and the Death of Copyright, 4 FIRST MONDAY (1999), <https://firstmonday.org/ojs/index.php/fm/article/view/684/594>.

[Source 154] Id.

[Source 155] Apache License, Version 2.0 § 7 (2004).

[Source 156] RESTATEMENT (THIRD) OF TORTS, supra note 128.

[Source 157] Id.

[Source 158] For Export Controls on AI, supra note 109.

[Source 159] Id.

[Source 160] DATACAMP, supra note 113.

[Source 161] Id.

[Source 162] Id.

WATERMARKING (ADDITIONAL)

[Source 163] DATACAMP, supra note 113.

[Source 164] Id.

[Source 165] Id.

[Source 166] Toward Reliable Provenance in AI-Generated Content: Text, Images, and Code, MEDIUM (Adnan Masood), <https://medium.com/@adnanmasood/toward-reliable-provenance-in-ai-generated-content-text-images-and-code-9ebe8c57ceae>.

[Source 167] DATACAMP, supra note 113.

[Source 168] TECHTARGET, supra note 115.

[Source 169] BROOKINGS, supra note 116.

[Source 170] ITU, supra note 117.

[Source 171] Id.

[Source 172] BROOKINGS, supra note 116.

[Source 173] EUROPEAN PARLIAMENT, supra note 120.

LIABILITY THEORY

[Source 174] Smith et al., supra note 23.

[Source 175] Id.

[Source 176] Id.

[Source 177] Id.

[Source 178] RESTATEMENT (THIRD) OF TORTS, supra note 128.

[Source 179] Id.

[Source 180] Calabresi, supra note 133.

[Source 181] RESTATEMENT (THIRD) OF TORTS: PRODUCTS LIABILITY, supra note 134.

[Source 182] 42 U.S.C. § 2210, supra note 135.

[Source 183] Calabresi, supra note 133.

[Source 184] Hu et al., supra note 11.

[Source 185] Id.

[Source 186] TECHSTRONG.AI, supra note 112.

[Source 187] RESTATEMENT (THIRD) OF AGENCY § 2.04 (AM. LAW INST. 2006).

[Source 188] Id.

[Source 189] Protocol of 2003 to the International Convention on the Establishment of an International Fund for Compensation for Oil Pollution Damage, May 16, 2003, IMO Doc. LEG/CONF.14/20.

[Source 190] Id.

[Source 191] Id.

[Source 192] Id.

[Source 193] Id.

[Source 194] Id.

NOTE ON CITATION FORMAT

All citations follow The Bluebook: A Uniform System of Citation (21st ed. 2020). Citations include:

- **Case law:** Full case name, reporter volume, reporter abbreviation, page number, court and year in parentheses
- **Statutes:** Title, codification, section, and year
- **Restatements:** Full Restatement title, section, and American Law Institute year
- **Books:** Author name, TITLE IN SMALL CAPS, page numbers (year)
- **Law reviews:** Author, Title, Volume JOURNAL ABBREVIATION page number (year)
- **Web sources:** Title, WEBSITE IN SMALL CAPS (date), URL
- **arXiv preprints:** Author et al., Title, arXiv:paper_number (year)
- **International treaties:** Treaty name, date, treaty series or document number

For subsequent references to the same source, "Id." or "supra note X" is used as appropriate per Bluebook rules.